

Adaptive Experimentation Methodology Note

This prototype is grounded in contextual bandits, adaptive experimentation, and OPE-style estimation. It is designed as decision support, not causal proof.

The demo uses simulated logged campaign interactions with assignment probabilities to estimate the effect of different actions. Estimates are most useful when the historical policy explored similar actions with adequate